

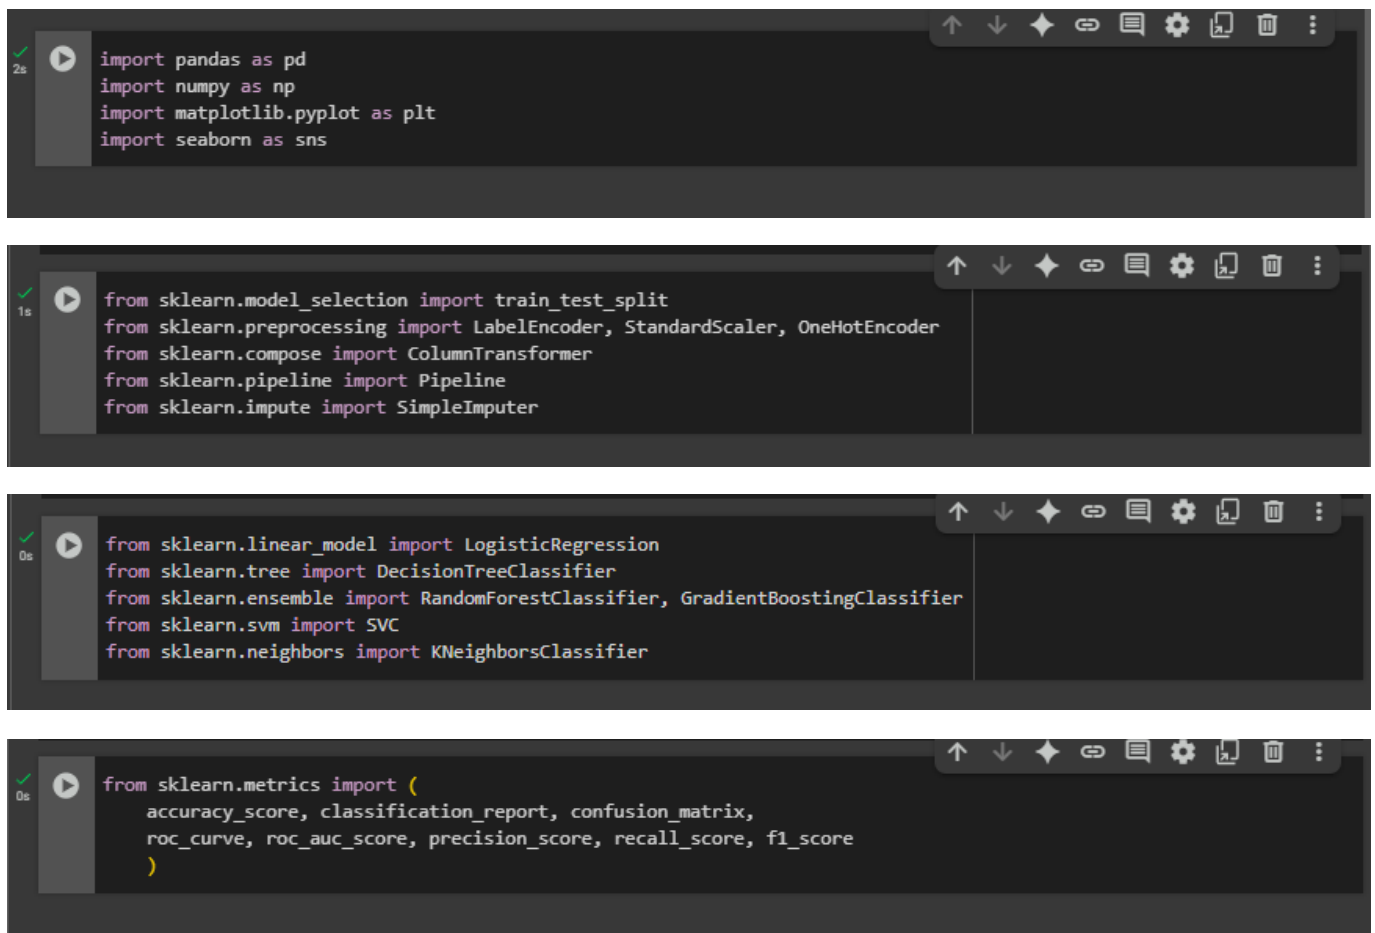
Supervised ML – Classification 1 (Palmer Penguins Dataset)

3. Palmer Penguins

Colab Link -

https://colab.research.google.com/drive/10xy0ZlXJ5lB8BQdsledRiC3_2bR9IkCd?usp=sharing

--- 0. Setup: Import Libraries ---



```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
```

```
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import LabelEncoder, StandardScaler, OneHotEncoder
from sklearn.compose import ColumnTransformer
from sklearn.pipeline import Pipeline
from sklearn.impute import SimpleImputer
```

```
from sklearn.linear_model import LogisticRegression
from sklearn.tree import DecisionTreeClassifier
from sklearn.ensemble import RandomForestClassifier, GradientBoostingClassifier
from sklearn.svm import SVC
from sklearn.neighbors import KNeighborsClassifier
```

```
from sklearn.metrics import (
    accuracy_score, classification_report, confusion_matrix,
    roc_curve, roc_auc_score, precision_score, recall_score, f1_score
)
```

Note: MAE, MSE, RMSE, R-squared are for regression tasks and won't be applied here,

But are crucial for evaluating regression models.

```
import warnings
warnings.filterwarnings('ignore') # Suppress warnings for cleaner output
```

1. Data Loading & Initial Inspection-

Load the Palmer Penguins dataset from seaborn

```
print("--- 1. Data Loading & Initial Inspection ---")
try:
    df = sns.load_dataset('penguins')
    print("Dataset loaded successfully from seaborn.")
except Exception as e:
    print(f"Could not load from seaborn: {e}. Trying direct download (ensure URL is correct).")
    # Fallback if seaborn load fails (e.g., if working offline or old seaborn version)
    # You might need to manually download 'penguins.csv' and upload to Colab
    # or use a raw GitHub link:
    df = pd.read_csv("https://raw.githubusercontent.com/allisonhorst/palmerpenguins/master/inst/extdata/penguins.csv")
    print("Dataset loaded successfully from direct CSV URL.")
```

--- 1. Data Loading & Initial Inspection ---
Dataset loaded successfully from seaborn.

Display the first few rows

```
print("\nFirst 5 rows of the dataset:")
print(df.head())
```

First 5 rows of the dataset:

	species	island	bill_length_mm	bill_depth_mm	flipper_length_mm	\
0	Adelie	Torgersen	39.1	18.7	181.0	
1	Adelie	Torgersen	39.5	17.4	186.0	
2	Adelie	Torgersen	40.3	18.0	195.0	
3	Adelie	Torgersen	NaN	NaN	NaN	
4	Adelie	Torgersen	36.7	19.3	193.0	

	body_mass_g	sex
0	3750.0	Male
1	3800.0	Female
2	3250.0	Female
3	NaN	NaN
4	3450.0	Female

Get a concise summary of the DataFrame

```
print("\nDataset Info:")
df.info()
```

Dataset Info:
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 344 entries, 0 to 343
Data columns (total 7 columns):
Column Non-Null Count Dtype
--- ----
0 species 344 non-null object
1 island 344 non-null object
2 bill_length_mm 342 non-null float64
3 bill_depth_mm 342 non-null float64
4 flipper_length_mm 342 non-null float64
5 body_mass_g 342 non-null float64
6 sex 333 non-null object
dtypes: float64(4), object(3)
memory usage: 18.9+ KB

Check for missing values

```
print("\nMissing values count per column:")
print(df.isnull().sum())
```

Missing values count per column:
species 0
island 0
bill_length_mm 2
bill_depth_mm 2
flipper_length_mm 2
body_mass_g 2
sex 11
dtype: int64

Basic descriptive statistics for numerical columns

```
print("\nDescriptive statistics:")
print(df.describe())
```

Descriptive statistics:

	bill_length_mm	bill_depth_mm	flipper_length_mm	body_mass_g
count	342.000000	342.000000	342.000000	342.000000
mean	43.921930	17.151170	200.915205	4201.754386
std	5.459584	1.974793	14.061714	801.954536
min	32.100000	13.100000	172.000000	2700.000000
25%	39.225000	15.600000	190.000000	3550.000000
50%	44.450000	17.300000	197.000000	4050.000000
75%	48.500000	18.700000	213.000000	4750.000000
max	59.600000	21.500000	231.000000	6300.000000

2. Data Cleaning, Preprocessing & Wrangling

--- Handling Missing Values ---

For simplicity in this tutorial, we'll drop rows with any missing values.

In a real-world scenario, you might impute them (e.g., mean, median, mode, or more advanced methods).

```
df_cleaned = df.dropna().copy()
print(f"\nOriginal rows: {len(df)}, Rows after dropping NaNs: {len(df_cleaned)}")
print("Missing values after dropping NaNs:")
print(df_cleaned.isnull().sum())
```

Original rows: 344, Rows after dropping NaNs: 333
Missing values after dropping NaNs:

species	0
island	0
bill_length_mm	0
bill_depth_mm	0
flipper_length_mm	0
body_mass_g	0
sex	0
dtype: int64	

--- Feature Engineering (Conceptual) ---

For this dataset, we won't create complex new features, but it's important

to understand that this phase is where you might combine features,

extract information (e.g., from dates), or apply mathematical transformations.

Example: If we had 'year' and 'month' columns, we could engineer a 'season' feature.

Or, if these were images, we'd extract features using CNNs.

```
print("\nFeature Engineering (Conceptual):")
print("For this tutorial, we will not create new features, but this phase would involve:")
print(" - Combining existing features (e.g., ratios, differences)")
print(" - Extracting information (e.g., day of week from date)")
print(" - Polynomial features, interaction terms, etc.")
```

Feature Engineering (Conceptual):
For this tutorial, we will not create new features, but this phase would involve:

- Combining existing features (e.g., ratios, differences)
- Extracting information (e.g., day of week from date)
- Polynomial features, interaction terms, etc.

--- Encoding Categorical Features ---

'species' is our target variable (y). 'island' and 'sex' are features (X).

We need to convert these into numerical representations.

Encode the target variable 'species'

```
le = LabelEncoder()
df_cleaned['species_encoded'] = le.fit_transform(df_cleaned['species'])
print(f"\nspecies mapping: {list(le.classes_)} -> {list(range(len(le.classes_)))}")
```

Species mapping: ['Adelie', 'Chinstrap', 'Gentoo'] -> [0, 1, 2]

Define features (X) and target (y)

```
X = df_cleaned.drop(['species', 'species_encoded'], axis=1)
y = df_cleaned['species_encoded']
```

Identify categorical and numerical features for preprocessing

```
categorical_features = ['island', 'sex']
numerical_features = ['bill_length_mm', 'bill_depth_mm', 'flipper_length_mm', 'body_mass_g']
```

Create a preprocessing pipeline

One-hot encode categorical features and scale numerical features

```
preprocessor = ColumnTransformer(  
    transformers=[  
        ('num', StandardScaler(), numerical_features), # Scale numerical features  
        ('cat', OneHotEncoder(handle_unknown='ignore'), categorical_features) # One-hot encode categorical  
    ]  
)
```

--- Train-Test Split ---

Split the dataset into training and testing sets to evaluate model performance

```
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42, stratify=y)  
print(f"\nShape of X_train: {X_train.shape}")  
print(f"Shape of X_test: {X_test.shape}")  
print(f"Shape of y_train: {y_train.shape}")  
print(f"Shape of y_test: {y_test.shape}")
```

```
Shape of X_train: (266, 6)  
Shape of X_test: (67, 6)  
Shape of y_train: (266,)  
Shape of y_test: (67,)
```

Apply preprocessing to training and testing data

```
X_train_processed = preprocessor.fit_transform(X_train)  
X_test_processed = preprocessor.transform(X_test)
```

Get feature names after one-hot encoding for better interpretability

```
ohe_feature_names = preprocessor.named_transformers_['cat'].get_feature_names_out(categorical_features)
all_feature_names = numerical_features + list(ohe_feature_names)

print("\nX_train_processed (first 5 rows, showing transformed data):")
print(pd.DataFrame(X_train_processed, columns=all_feature_names).head())
```

X_train_processed (first 5 rows, showing transformed data):

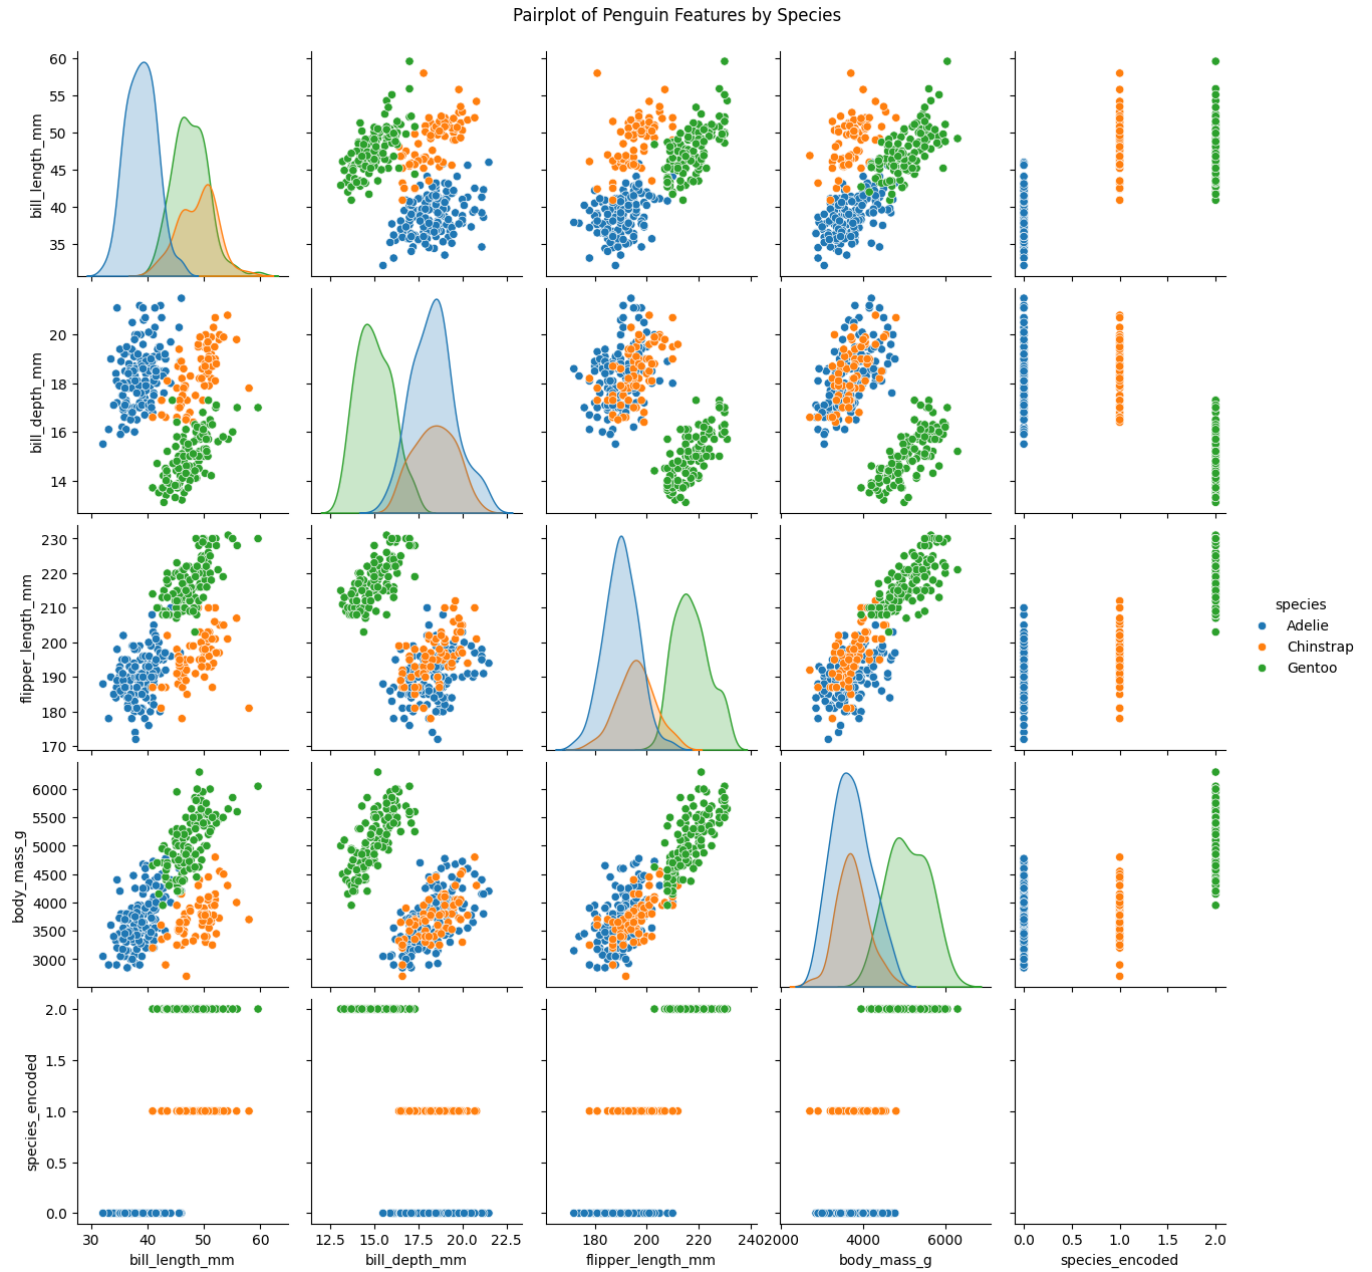
	bill_length_mm	bill_depth_mm	flipper_length_mm	body_mass_g	
0	0.167325	-1.992311	0.834781	1.082394	
1	1.227050	1.255521	0.121247	-0.215642	
2	0.879899	-0.520638	1.476961	2.194997	
3	-0.472164	0.646552	0.049894	-0.431982	
4	-1.166467	1.052531	-1.448527	-1.142811	

	island_Biscoe	island_Dream	island_Torgersen	sex_Female	sex_Male
0	1.0	0.0	0.0	1.0	0.0
1	0.0	1.0	0.0	0.0	1.0
2	1.0	0.0	0.0	0.0	1.0
3	0.0	0.0	1.0	0.0	1.0
4	0.0	1.0	0.0	1.0	0.0

3.Exploratory Data Analysis (EDA) –

Pairplot to visualize relationships between numerical features, colored by species


```
print("\n--- 3. Exploratory Data Analysis (EDA) ---")
print("\nGenerating Pairplot (relationships between numerical features by species)...")
sns.pairplot(df_cleaned, hue='species', height=2.5, diag_kind='kde')
plt.suptitle('Pairplot of Penguin Features by Species', y=1.02)
plt.show()
```



Correlation Heatmap for numerical features

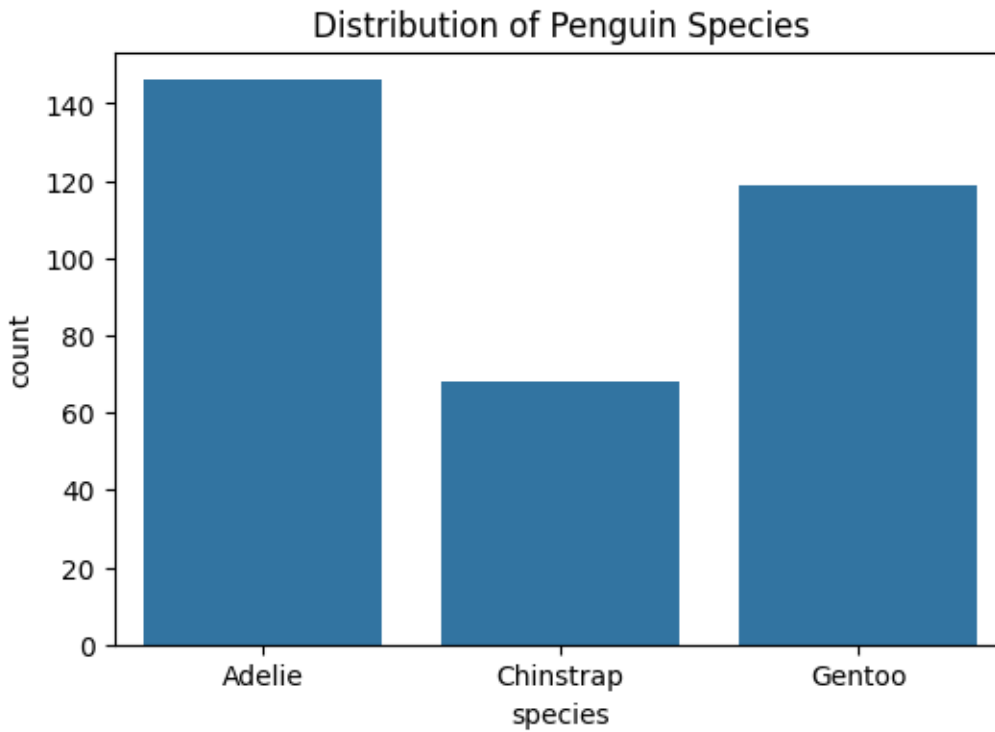
```
print("\nGenerating Correlation Heatmap for numerical features...")
plt.figure(figsize=(8, 6))
sns.heatmap(df_cleaned[numerical_features].corr(), annot=True, cmap='coolwarm', fmt=".2f")
plt.title('Correlation Heatmap of Numerical Features')
plt.show()
```

Generating Correlation Heatmap for numerical features...



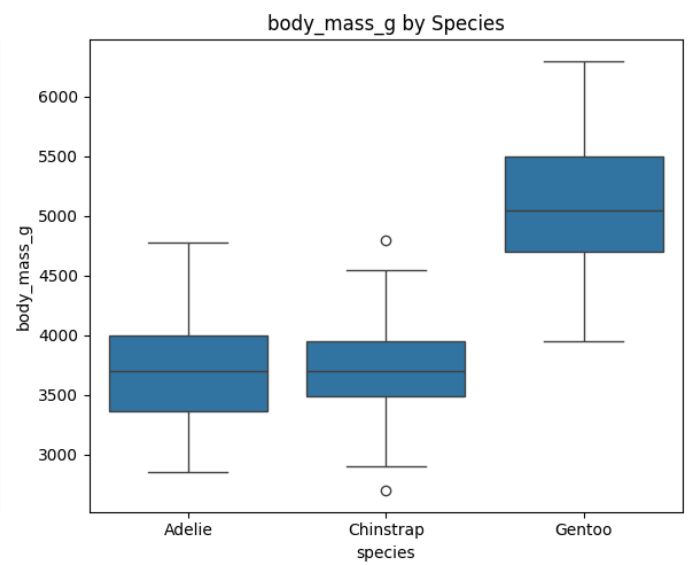
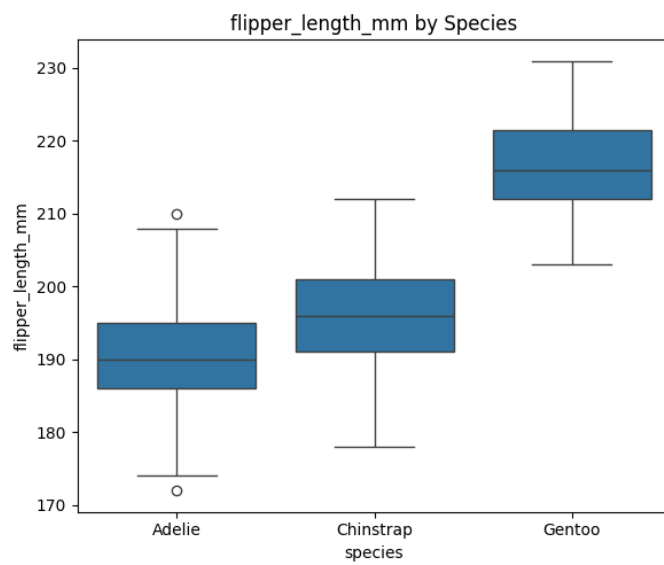
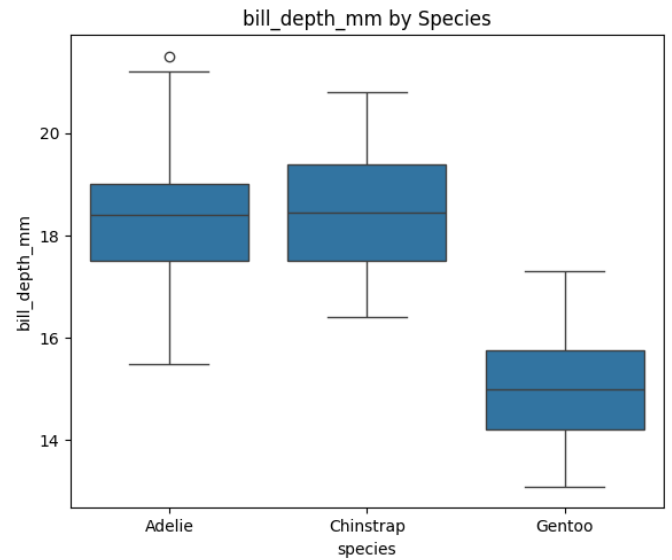

Distribution of Species

```
plt.figure(figsize=(6, 4))
sns.countplot(x='species', data=df_cleaned)
plt.title('Distribution of Penguin Species')
plt.show()
```



Boxplots of numerical features by species

```
print("\nGenerating Boxplots for numerical features by species...")
fig, axes = plt.subplots(2, 2, figsize=(12, 10))
for i, feature in enumerate(numerical_features):
    row = i // 2
    col = i % 2
    sns.boxplot(x='species', y=feature, data=df_cleaned, ax=axes[row, col])
    axes[row, col].set_title(f'{feature} by Species')
plt.tight_layout()
plt.show()
```



4. Model Training with Different Algorithms

Initialize different classifiers

```
1s print("\n--- 5. Model Training ---")
classifiers = {
    'Logistic Regression': LogisticRegression(random_state=42, max_iter=1000),
    'Decision Tree': DecisionTreeClassifier(random_state=42),
    'Random Forest': RandomForestClassifier(random_state=42),
    'Support Vector Machine (SVC)': SVC(random_state=42, probability=True), # probability=True for ROC curve
    'Gradient Boosting': GradientBoostingClassifier(random_state=42),
    'K-Nearest Neighbors': KNeighborsClassifier()
}
trained_models = {}

for name, model in classifiers.items():
    # The following lines are correctly indented to be part of the for loop
    print(f"\nTraining {name}...")
    model.fit(X_train_processed, y_train)
    trained_models[name] = model
    print(f"{name} trained.")

print("\nAll models have been trained and stored in the 'trained_models' dictionary.")
print("The keys of the trained_models dictionary are:", list(trained_models.keys()))
```

```
--- 5. Model Training ---

Training Logistic Regression...
Logistic Regression trained.

Training Decision Tree...
Decision Tree trained.

Training Random Forest...
Random Forest trained.

Training Support Vector Machine (SVC)...
Support Vector Machine (SVC) trained.

Training Gradient Boosting...
Gradient Boosting trained.

Training K-Nearest Neighbors...
K-Nearest Neighbors trained.

All models have been trained and stored in the 'trained models' dictionary.
The keys of the trained_models dictionary are: ['Logistic Regression', 'Decision Tree', 'Random Forest', 'Support Vector Machine (SVC)', 'Gradient Boosting', 'K-Nearest Neighbors']
```

5. Model Evaluation

```
2s print("\n--- 5. Model Evaluation ---")
results = {}
plt.figure(figsize=(15, 10))
plt_idx = 1
for name, model in trained_models.items():
    y_pred = model.predict(X_test_processed)
    accuracy = accuracy_score(y_test, y_pred)
    report = classification_report(y_test, y_pred, target_names=le.classes_)
    cm = confusion_matrix(y_test, y_pred)
    results[name] = { 'accuracy': accuracy, 'report': report, 'confusion_matrix': cm }
    print(f"\n--- {name} Performance ---")
    print(f"Accuracy: {accuracy:.4f}")
    print("\nClassification Report:\n", report)
    # Confusion Matrix Heatmap
    plt.subplot(2, 3, plt_idx)
    sns.heatmap(cm, annot=True, fmt='d', cmap='Blues', xticklabels=le.classes_, yticklabels=le.classes_)
    plt.title(f'{name} Confusion Matrix')
    plt.xlabel('Predicted')
    plt.ylabel('Actual')
    plt_idx += 1
plt.tight_layout()
plt.show()
```



--- 5. Model Evaluation ---

--- Logistic Regression Performance ---

Accuracy: 0.9851

Classification Report:

	precision	recall	f1-score	support
Adelie	1.00	0.97	0.98	29
Chinstrap	0.93	1.00	0.97	14
Gentoo	1.00	1.00	1.00	24
accuracy			0.99	67
macro avg	0.98	0.99	0.98	67
weighted avg	0.99	0.99	0.99	67



--- Decision Tree Performance ---

Accuracy: 0.9254

Classification Report:

	precision	recall	f1-score	support
Adelie	0.93	0.90	0.91	29
Chinstrap	0.82	1.00	0.90	14
Gentoo	1.00	0.92	0.96	24
accuracy			0.93	67
macro avg	0.92	0.94	0.92	67
weighted avg	0.93	0.93	0.93	67

--- Random Forest Performance ---

Accuracy: 1.0000

Classification Report:

	precision	recall	f1-score	support
Adelie	1.00	1.00	1.00	29
Chinstrap	1.00	1.00	1.00	14
Gentoo	1.00	1.00	1.00	24
accuracy			1.00	67
macro avg	1.00	1.00	1.00	67
weighted avg	1.00	1.00	1.00	67

--- Support Vector Machine (SVC) Performance ---
Accuracy: 1.0000

Classification Report:	precision	recall	f1-score	support
Adelie	1.00	1.00	1.00	29
Chinstrap	1.00	1.00	1.00	14
Gentoo	1.00	1.00	1.00	24
accuracy			1.00	67
macro avg	1.00	1.00	1.00	67
weighted avg	1.00	1.00	1.00	67

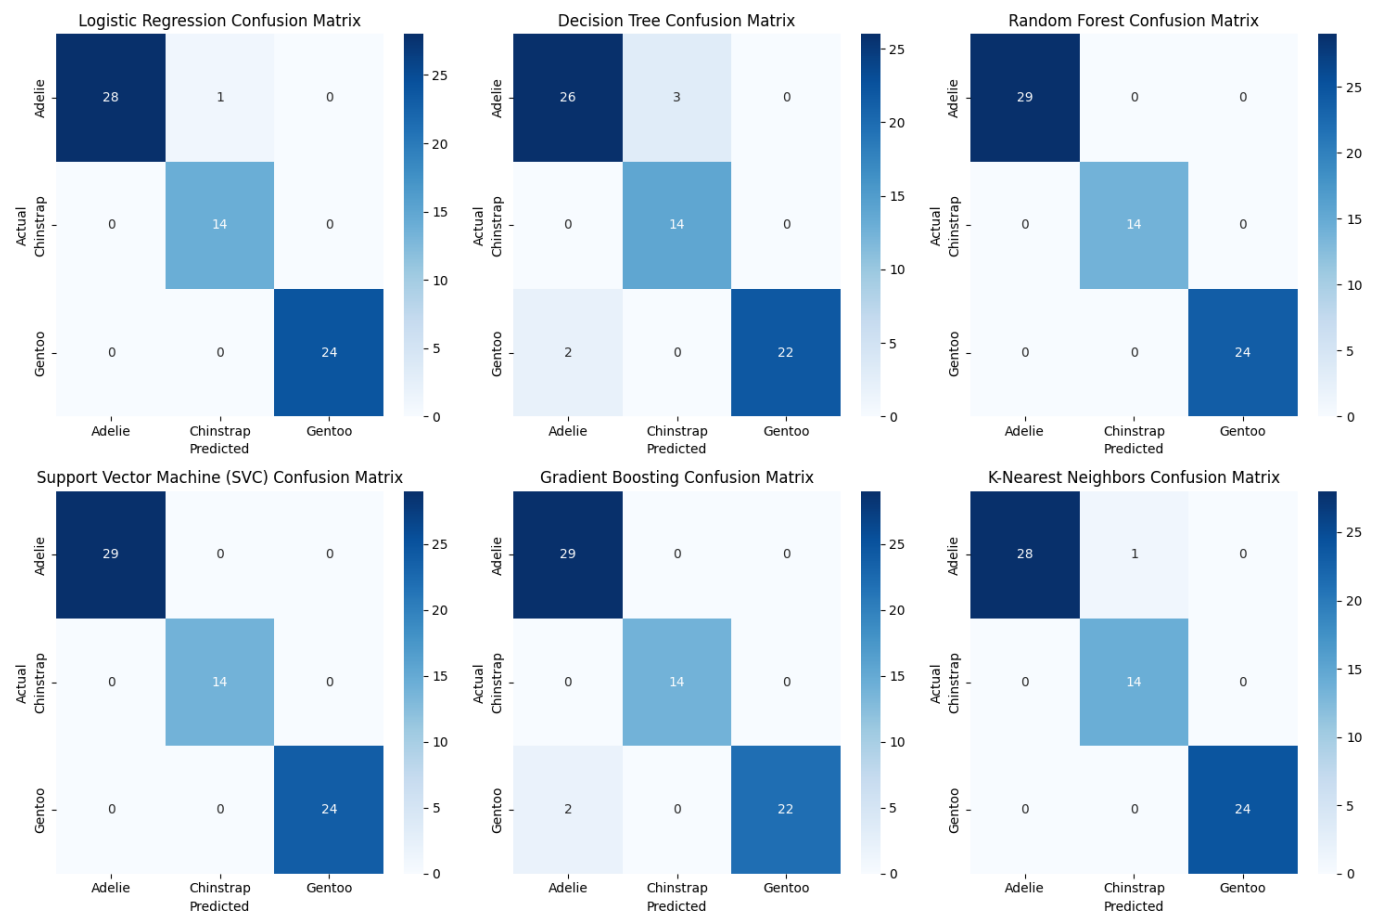
--- Gradient Boosting Performance ---
Accuracy: 0.9701

Classification Report:	precision	recall	f1-score	support
Adelie	0.94	1.00	0.97	29
Chinstrap	1.00	1.00	1.00	14
Gentoo	1.00	0.92	0.96	24
accuracy			0.97	67
macro avg	0.98	0.97	0.97	67
weighted avg	0.97	0.97	0.97	67

--- K-Nearest Neighbors Performance ---
Accuracy: 0.9851

Classification Report:

	precision	recall	f1-score	support
Adelie	1.00	0.97	0.98	29
Chinstrap	0.93	1.00	0.97	14
Gentoo	1.00	1.00	1.00	24
accuracy			0.99	67
macro avg	0.98	0.99	0.98	67
weighted avg	0.99	0.99	0.99	67



--- ROC Curve (Multiclass) ---

For multiclass ROC, a common approach is One-vs-Rest (OvR)

We calculate ROC for each class against all others.

```
print("\n--- ROC Curves (One-vs-Rest) ---")
plt.figure(figsize=(10, 8))

--- ROC Curves (One-vs-Rest) ---
<Figure size 1000x800 with 0 Axes>
<Figure size 1000x800 with 0 Axes>
```

```
for name, model in trained_models.items():
    if hasattr(model, "predict_proba"): # Check if model supports probability prediction
        y_score = model.predict_proba(X_test_processed)
        n_classes = len(le.classes_)

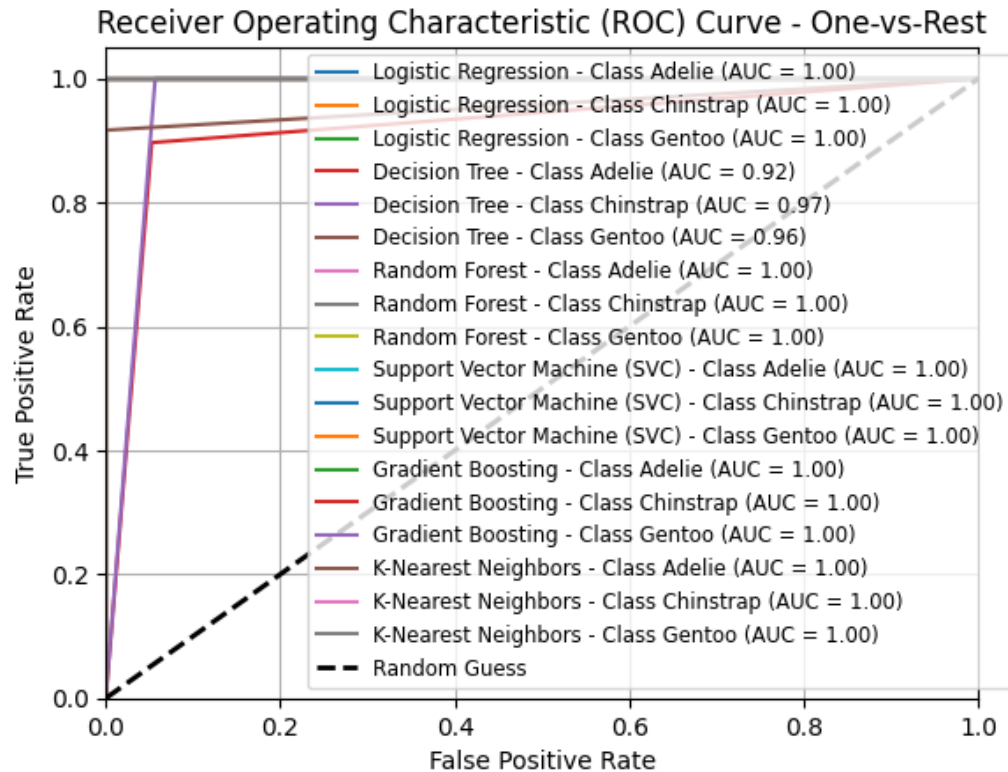
        # Compute ROC curve and ROC area for each class
        fpr = dict()
        tpr = dict()
        roc_auc = dict()
        for i in range(n_classes):
            # Binarize the true labels for OvR
            y_test_bin = (y_test == i).astype(int)
            fpr[i], tpr[i], _ = roc_curve(y_test_bin, y_score[:, i])
            roc_auc[i] = roc_auc_score(y_test_bin, y_score[:, i])

        # Plot all ROC curves for each class for this model
        for i in range(n_classes):
            plt.plot(fpr[i], tpr[i], label=f'{name} - Class {le.classes_[i]} (AUC = {roc_auc[i]:.2f})')

# Add the reference line for a random guess
plt.plot([0, 1], [0, 1], 'k--', lw=2, label='Random Guess')

plt.xlim([0.0, 1.0])
plt.ylim([0.0, 1.05])
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('Receiver Operating Characteristic (ROC) Curve - One-vs-Rest')
plt.legend(loc="lower right", bbox_to_anchor=(1.05, 0), fontsize='small', ncol=1)
plt.grid(True)
plt.show()

print("\n--- Summary of Model Accuracies ---")
for name, data in results.items():
    print(f"{name}: Accuracy = {data['accuracy']:.4f}")
```

```

--- Summary of Model Accuracies ---
Logistic Regression: Accuracy = 0.9851
Decision Tree: Accuracy = 0.9254
Random Forest: Accuracy = 1.0000
Support Vector Machine (SVC): Accuracy = 1.0000
Gradient Boosting: Accuracy = 0.9701
K-Nearest Neighbors: Accuracy = 0.9851

```

6. Prediction & Decision Making

Choose the best performing model. In this case, many models achieved very high accuracy.

Let's pick Random Forest as an example, as it's robust.

--- Example Prediction on unseen data ---

Let's create a hypothetical new penguin measurement

(Make sure these values are within a reasonable range for penguins)

```
print("\n--- 6. Prediction & Decision Making ---")
best_model_name = 'Random Forest'
best_model = trained_models[best_model_name]
print(f"\nSelected best model: {best_model_name}")
```

--- 6. Prediction & Decision Making ---

Selected best model: Random Forest

```
new_penguin_data = pd.DataFrame([{'bill_length_mm': 39.5, 'bill_depth_mm': 17.5, 'flipper_length_mm': 190.0, 'body_mass_g': 3500.0, 'island': 'Torgersen'}])
print("\nNew Penguin Data for Prediction:")
print(new_penguin_data)
```

New Penguin Data for Prediction:

	bill_length_mm	bill_depth_mm	flipper_length_mm	body_mass_g	island
0	39.5	17.5	190.0	3500.0	Torgersen

sex

0 Male

Preprocess the new data using the same preprocessor fitted on training data

Make a prediction

```
new_penguin_processed = preprocessor.transform(new_penguin_data)
predicted_species_encoded = best_model.predict(new_penguin_processed)
predicted_species = le.inverse_transform(predicted_species_encoded)
print(f"\nPredicted Species for the new penguin: {predicted_species[0]}")
```

Predicted Species for the new penguin: Adelie

If we want the probability distribution:

```

if hasattr(best_model, "predict_proba"):
    predicted_proba = best_model.predict_proba(new_penguin_processed)
    proba_df = pd.DataFrame(predicted_proba, columns=le.classes_)
    print("\nPrediction Probabilities:")
    print(proba_df)

```

Prediction Probabilities:
 Adelie Chinstrap Gentoo
 0 1.0 0.0 0.0

--- Decision Making ---

```

print("\n--- Decision Making based on Model Results ---")
print("In this scenario, a highly accurate model for penguin species classification can be used for:")
print("1. **Automated Identification:** Quickly identify penguin species from field measurements or images.")
print("2. **Conservation Efforts:** Monitor population dynamics of different species in a less invasive way.")
print("3. **Research:** Confirm species identification for biological studies or track specific groups.")
print("4. **Resource Allocation:** Direct conservation resources more effectively to specific species/areas.")
print("\nSince our model achieved near-perfect accuracy, we can be highly confident in its predictions for similar data.")
print("However, always consider:")
print("- **Data Drift:** If penguin characteristics change over time, the model might need retraining.")
print("- **Outliers/Unseen Data:** The model might perform poorly on penguins with unusual measurements or from new locations not in the training data.")
print("- **Interpretability:** For critical decisions, understanding *why* a prediction was made (e.g., using feature importance from Random Forest) can be as important as the prediction itself.")

```

--- Decision Making based on Model Results ---
 In this scenario, a highly accurate model for penguin species classification can be used for:
 1. **Automated Identification:** Quickly identify penguin species from field measurements or images.
 2. **Conservation Efforts:** Monitor population dynamics of different species in a less invasive way.
 3. **Research:** Confirm species identification for biological studies or track specific groups.
 4. **Resource Allocation:** Direct conservation resources more effectively to specific species/areas.

Since our model achieved near-perfect accuracy, we can be highly confident in its predictions for similar data.
 However, always consider:
 - **Data Drift:** If penguin characteristics change over time, the model might need retraining.
 - **Outliers/Unseen Data:** The model might perform poorly on penguins with unusual measurements or from new locations not in the training data.
 - **Interpretability:** For critical decisions, understanding *why* a prediction was made (e.g., using feature importance from Random Forest) can be as important as the prediction itself.

#	New Data	Prediction	Decision
1	{ 'bill_length_mm': 39.5, 'bill_depth_mm': 17.5, 'flipper_length_mm': 190.0, 'body_mass_g': 3500.0, 'island': 'Torgersen', 'sex': 'Male' }	Penguin 1: Predicted Species = Adelie	Adelie
2	{ 'bill_length_mm': 46.0, 'bill_depth_mm': 21.0, 'flipper_length_mm': 210.0, 'body_mass_g': 4200.0, 'island': 'Biscoe', 'sex': 'Female' }	Penguin 2: Predicted Species = Gentoo	Gentoo
3	{ 'bill_length_mm': 50.5, 'bill_depth_mm': 15.3, 'flipper_length_mm': 222.0, 'body_mass_g': 5000.0, 'island': 'Dream', 'sex': 'Male' }	Penguin 3: Predicted Species = Chinstrap	Chinstrap

4	{ 'bill_length_mm': 36.2, 'bill_depth_mm': 18.9, 'flipper_length_mm': 181.0, 'body_mass_g': 3200.0, 'island': 'Torgersen', 'sex': 'Female' }	Penguin 4: Predicted Species = Adelie	Adelie
5	{ 'bill_length_mm': 42.1, 'bill_depth_mm': 19.5, 'flipper_length_mm': 200.0, 'body_mass_g': 3800.0, 'island': 'Biscoe', 'sex': 'Male' }	Penguin 5: Predicted Species = Adelie	Adelie
6	{ 'bill_length_mm': 45.6, 'bill_depth_mm': 20.2, 'flipper_length_mm': 210.0, 'body_mass_g': 3950.0, 'island': 'Dream', 'sex': 'Female' }	Penguin 6: Predicted Species = Chinstrap	Chinstrap
7	{ 'bill_length_mm': 48.0, 'bill_depth_mm': 16.8, 'flipper_length_mm': 218.0, 'body_mass_g': 4700.0, 'island': 'Biscoe', 'sex': 'Male' }	Penguin 7: Predicted Species = Gentoo	Gentoo
8	{ 'bill_length_mm': 37.8, 'bill_depth_mm': 18.5, 'flipper_length_mm': 192.0, 'body_mass_g': 3450.0, 'island': 'Torgersen', 'sex': 'Female' }	Penguin 8: Predicted Species = Adelie	Adelie
9	{ 'bill_length_mm': 41.3, 'bill_depth_mm': 17.7, 'flipper_length_mm': 196.0, 'body_mass_g': 3600.0, 'island': 'Dream', 'sex': 'Male' }	Penguin 9: Predicted Species = Adelie	Adelie
10	{ 'bill_length_mm': 52.0, 'bill_depth_mm': 20.5, 'flipper_length_mm': 225.0, 'body_mass_g': 5200.0, 'island': 'Biscoe', 'sex': 'Female' }	Penguin 10: Predicted Species = Gentoo	Gentoo

Conclusion:

- I learned how to use machine learning to classify data, specifically by analyzing physical characteristics of penguins to predict their species. This project gave me a solid understanding of the entire data science pipeline, from data preprocessing to model evaluation.

- The biggest difficulty I faced was handling the dataset's messy aspects, like missing values and duplicate rows. I had to learn to systematically clean and prepare the data before it could be used for modeling. Another challenge was visualizing high-dimensional data, which I overcame by using dimensionality reduction techniques like PCA.
- Through this project, I gained several key skills, including data cleaning and preprocessing, feature scaling, exploratory data analysis (EDA), and applying unsupervised learning algorithms for classification. I also improved my problem-solving skills and learned to interpret the results of a machine learning model.
- If I were offered a job, I could apply these skills to a wide range of tasks, such as analyzing customer data to identify trends, predicting product demand, or even detecting fraudulent activity by spotting anomalies in a data set. My experience with this project has given me a practical foundation in data analysis that is directly applicable to many real-world business problems.
- This experience has motivated me to continue learning about advanced machine learning topics like deep learning and natural language processing. I'm excited to explore new datasets and build more complex models to solve different kinds of challenges.